

Function sbr

Description

Function for scalable Bayesian regression (SBR) and sparse SBR (SSBR) in normal linear models with multiple types (sources) of feature matrices (with K being the number of sources). When $K = 1$, SBR corresponds to standard ridge regression using one from the three available empirical Bayes estimators (see below) for the penalty parameter. For details see Perrakis, Mukherjee and the Alzheimers Disease Neuroimaging Initiative. (2018).

Usage

```
sbr(y, X, trX, G, estimator = "MAP", sparsify = FALSE, relaxed = TRUE, sparse.control = 1,
p.threshold = 5000, cov.blocks = 1000, parallel = FALSE, cl, L.optim = 1e-04, U.optim =
1e+04)
```

Arguments

y	a standardized response vector.
X	a standardized feature matrix (if $K = 1$) or a list of standardized feature matrices (if $K > 1$).
trX	(optional) the transpose matrix of X (if $K = 1$) or a list of transpose matrices (if $K > 1$).
G	the inner-product Gram matrix (if $K = 1$) or a list containing the multiple Gram matrices (if $K > 1$).
estimator	the estimator used for tuning the shrinkage levels. Available estimates are leave-one-out cross-validation ("CV"), maximum marginal likelihood ("ML") and the maximum-a-posteriori value ("MAP", default).
sparsify	logical, if TRUE the SSBR solution is calculated, default option is FALSE.
relaxed	logical, if TRUE (default) the relaxed SSBR solution is calculated, if FALSE the general SSBR solution is calculated.
sparse.control	numerical value for controlling the effect of sample size (n) on the resulting SSBR solution when relaxed = TRUE. Default option is 1 (no control). A recommended option for sparser solutions is sparse.control = log(n).
p.threshold	used for block-matrix computation of the main diagonal of the covariance matrix when sparsify = TRUE and relaxed = TRUE. It will be triggered for any source-matrix whose number of columns is larger than p.threshold.
cov.blocks	argument corresponding to block size (not the number of blocks) when the block-matrix computation is triggered (see above). Default option is 1000, i.e. blocks of dimensionality 1000×1000 .

parallel	logical, if <code>parallel = TRUE</code> the calculation of variance components required for the relaxed SSBR solution is performed in parallel. Default is <code>FALSE</code> .
cl	the number of cores to use when <code>parallel = TRUE</code> . Must be provided by the user.
L.optim	lower bound for the optimization procedure used to tune the shrinkage levels, default is <code>1e-04</code> .
U.optim	upper bound for the optimization procedure used to tune the shrinkage levels, default is <code>1e+04</code> .

Value

coefficients	a 1-column matrix with the SBR beta estimates (when <code>sparsify = FALSE</code>) or a 2-column matrix with the SBR and SSBR beta estimates (when <code>sparsify = TRUE</code>). Note that the coefficients correspond to the standardized response variable and feature matrix.
sigma2	the variance component (at the posterior mode).
lambda	the vector of penalty parameters.
lambdaEstimator	the estimator used for <code>lambda</code> .
duration	reported runtime.

References

Perrakis, K., Mukherjee, S. and the Alzheimers Disease Neuroimaging Initiative. (2018). *Scalable Bayesian regression in high dimensions with multiple data sources*, arXiv:1710.00596 [stat.ME].

Examples:

```
##### Example with 3 data sources #####

library(mvtnorm)
library(MCMCpack)

### GENERATION OF DATA ###

## sample size and number of predictors
n <- 50
p1 <- 10
p2 <- 100
p3 <- 300

## generation of covariance matrices and feature matrices
S1 <- riwish(p1, diag(p1))
S2 <- riwish(p2, diag(p2))
S3 <- riwish(p3, diag(p3))
X1 <- matrix(rmvnorm(n * p1, rep(0, p1), S1), n, p1)
X2 <- matrix(rmvnorm(n * p2, rep(0, p2), S2), n, p2)
X3 <- matrix(rmvnorm(n * p3, rep(0, p3), S3), n, p3)
```

```

## sparsity and generation of betas
s2 <- p2 * 0.3
s3 <- p3 * 0.01
non.zero2 <- sample(1:p2, s2)
non.zero3 <- sample(1:p3, s3)
b1 <- rnorm(10, 0, 2.5)
b2 <- rep(0, p2)
b2[non.zero2] <- rnorm(s2)
b3 <- rep(0, p3)
b3[non.zero3] <- rnorm(s3)

## generation of responce
mu <- X1 %*% b1 + X2 %*% b2 + X3 %*% b3
y <- rnorm(n, mu, sd=0.5)

## standardize
y <- scale(y)
X1 <- scale(X1)
X2 <- scale(X2)
X3 <- scale(X3)

## calculation of gram matrices
G1 <- X1 %*% t(X1); G2 <- X2 %*% t(X2); G3 <- X3 %*% t(X3)

## make lists
G <- list(G1, G2, G3)
X <- list(X1, X2, X3)

### RUN SBR/SSBR ###

# 1) SBR with the ML lambda-estimator
model1 <- sbr(y = y, X = X, G = G, estimator = 'ML')

# 2) relaxed SSBR with the ML lambda-estimator using block-matrix computations for the
# variances of X3 (since p3=300)
model2 <- sbr(y = y, X = X, G = G, estimator = 'ML', sparsify = TRUE, p.threshold =
  100, cov.blocks = 100)

# 3) SSBR with the ML lambda-estimator
model3 <- sbr(y = y, X = X, G = G, estimator = 'ML', sparsify = TRUE, relaxed = FALSE)

# 4) parallel computing for the configuration of model2
cores <- detectCores() - 1
cores <- makeCluster(cores)
registerDoParallel(cores)
model4 <- sbr(y = y, X = X, G = G, estimator = 'ML', parallel = TRUE, cl = cores,
  sparsify = TRUE, p.threshold = 100, cov.blocks = 100)
stopCluster(cores)

### EXTRACTING OUTPUT FROM A MODEL ###
coef(model3)      # SBR/SSBR coefficients (or alternatively model3$coefficients)
model3$lambda     # vector of lambdas
model3$sigma2     # error variance
model3$duration   # runtime

```

Function `gram`

Description

Function for calculating the (inner-product) Gram matrix that allows for block-matrix multiplication.

Usage

```
gram(X, trX, block = FALSE, block.size = 1000, show.progress = FALSE)
```

Arguments

<code>X</code>	a standardized feature matrix.
<code>trX</code>	(optional) the transpose matrix of <code>X</code> .
<code>block</code>	logical, block matrix computation is performed when <code>TRUE</code> , default option is <code>FALSE</code> .
<code>block.size</code>	used when <code>block = TRUE</code> . Default option is 1000, i.e. blocks of dimensionality 1000×1000 .
<code>show.progress</code>	logical, when <code>TRUE</code> (and <code>block = TRUE</code>) the progress of the calculations is reported on the console.

Value

Returns the inner-product Gram matrix.

Examples:

```
X <- matrix(rnorm(100 * 300), 100, 300)
G0 <- gram(X)                                # usual matrix multiplication
G1 <- gram(X, block=TRUE, block.size=100)    # block matrix multiplication
all.equal(G0, G1)
```

Function `gram.parallel`

Description

Function for calculating the (inner-product) Gram matrix that allows for block-matrix multiplication performed in parallel.

Usage

```
gram.parallel(X, cl, ...)
```

Arguments

- X** a standardized feature matrix.
- cl** the number of cores to use. Must be provided by the user.
- ... additional arguments passed from function `gram`. One can additionally use block matrix multiplication within each core. Warning: in this case argument `block.size` must not exceed $\text{floor}(\text{ncol}(X)/\text{number of cores})-1$.

Value

Returns the inner-product Gram matrix.

Examples:

```
X <- matrix(rnorm(100 * 300), 100, 300)
G0 <- gram(X)                                # usual matrix multiplication
cores <- detectCores() - 1
cores <- makeCluster(cores)
registerDoParallel(cores)
G1 <- gram.parallel(X, cl = cores)           # parallel matrix multiplication
stopCluster(cores)
all.equal(G0, G1)
```

Function `predict.sbr`

Description

Predict S3 method for objects of class 'sbr'.

Usage

```
predict.sbr(object, newdata, coef = 'sbr')
```

Arguments

- object** object an object of class 'sbr'.
- newdata** a (standardized) data matrix from which to predict.
- coef** choose whether to use the SBR beta estimates (default option "sbr") or the SSBR beta estimates (option "ssbr").

Value

Returns a vector of predictions.

Examples:

```
y <- rnorm(100)
X <- matrix(rnorm(100*300), 100, 300)
G <- gram(X)
model <- sbr(y = y, X = X, G = G, sparsify = TRUE)
predict1 <- predict(model, X[1:10, ], coef = 'sbr')
predict2 <- predict(model, X[1:10, ], coef = 'ssbr')
```